

NAPQUEENS ASSIGNMENT

20MIY0056

SUHAINA

```
In [18]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load the datasets
train = pd.read_csv('C:\\Users\\suhai\\Downloads\\train.csv\\train.csv')
test = pd.read_csv('C:\\Users\\suhai\\Downloads\\test.csv')
submission = pd.read_csv('C:\\Users\\suhai\\Downloads\\sample_submission.csv')

# Convert date columns to datetime
train['date'] = pd.to_datetime(train['date'])
test['date'] = pd.to_datetime(test['date'])
```

```
In [19]: # Basic information
print("Train dataset info:")
print(train.info())
print("\nTest dataset info:")
print(test.info())

print("\nTrain dataset statistics:")
print(train.describe())
print("\nTest dataset statistics:")
print(test.describe())

print("\nMissing values in train dataset:")
print(train.isnull().sum())
print("\nMissing values in test dataset:")
print(test.isnull().sum())

plt.figure(figsize=(10, 6))
sns.histplot(train['units'], bins=50, kde=True)
plt.title('Distribution of Units Sold')
plt.xlabel('Units Sold')
plt.ylabel('Frequency')
plt.show()
```

Train dataset info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101490 entries, 0 to 101489
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              101490 non-null object
1   date            101490 non-null datetime64[ns]
2   Item Id        101488 non-null object
3   Item Name      99658 non-null object
4   ad_spend       77303 non-null float64
5   anarix_id      101490 non-null object
6   units          83592 non-null float64
7   unit_price     101490 non-null float64
dtypes: datetime64[ns](1), float64(3), object(4)
memory usage: 6.2+ MB
None
```

Test dataset info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2833 entries, 0 to 2832
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              2833 non-null object
1   date            2833 non-null datetime64[ns]
2   Item Id        2833 non-null object
3   Item Name      2489 non-null object
4   ad_spend       1382 non-null float64
5   anarix_id      2833 non-null object
6   unit_price     2833 non-null float64
dtypes: datetime64[ns](1), float64(2), object(4)
memory usage: 155.1+ KB
None
```

Train dataset statistics:

	ad_spend	units	unit_price
count	77303.000000	83592.000000	101490.000000
mean	110.771470	10.284381	106.750922
std	529.303777	68.945915	425.704733
min	0.000000	-173.000000	-8232.000000
25%	0.000000	0.000000	0.000000
50%	4.230000	1.000000	0.000000
75%	44.310000	5.000000	0.000000
max	47934.990000	9004.000000	21557.390000

Test dataset statistics:

	ad_spend	unit_price
count	1382.000000	2833.000000
mean	198.838032	98.725873
std	797.354508	383.585307
min	0.000000	-1988.180000
25%	0.730000	0.000000
50%	39.200000	0.000000
75%	156.012500	0.000000
max	18724.850000	6870.000000

Missing values in train dataset:


ID	0
date	0
Item Id	2
Item Name	1832
ad_spend	24187
anarix_id	0
units	17898
unit_price	0


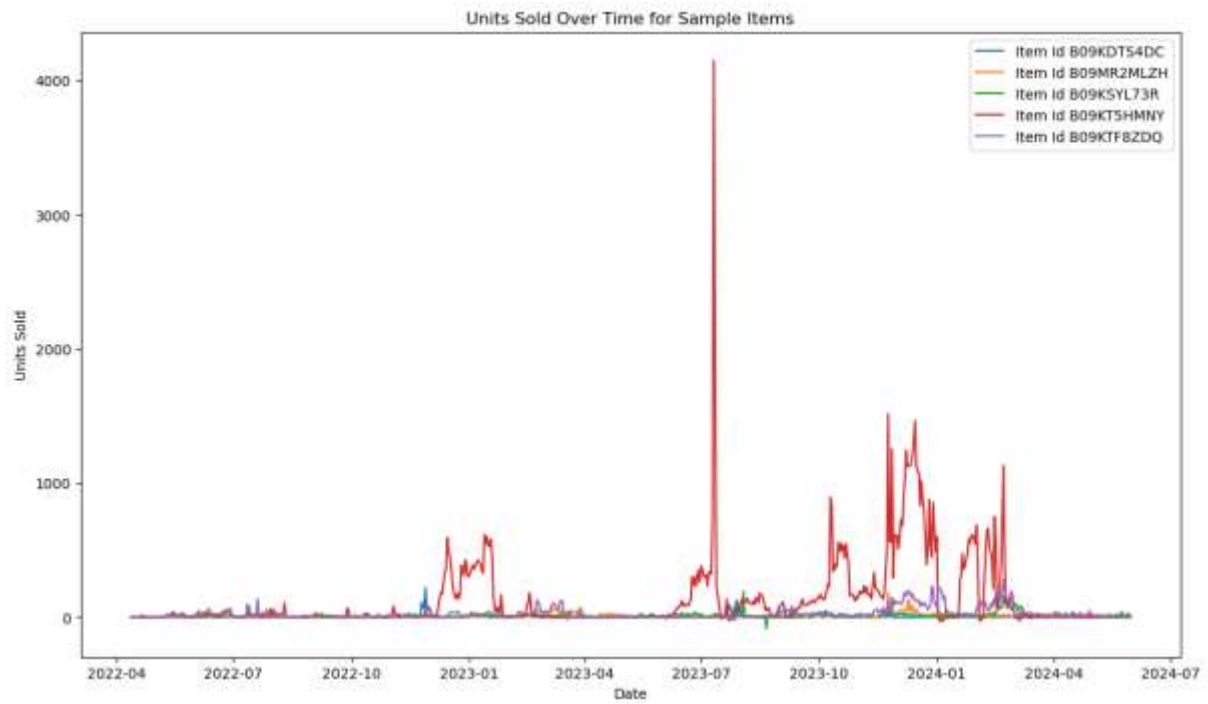
dtype: int64

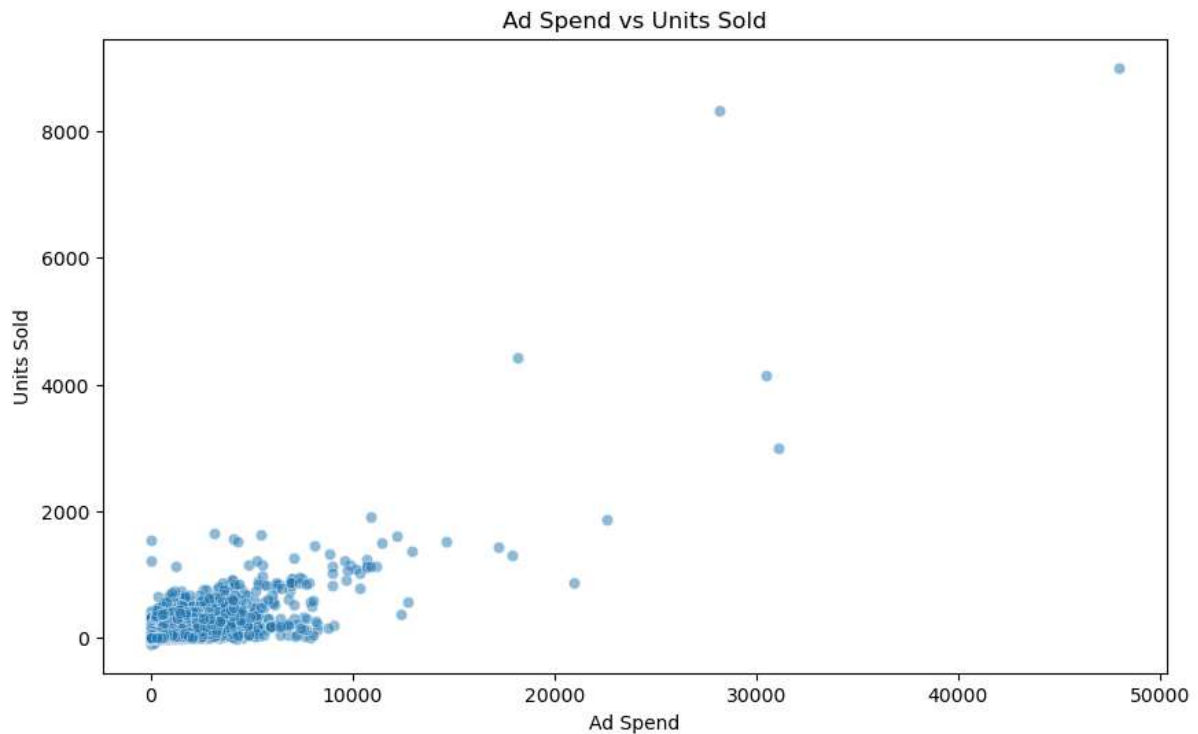
Missing values in test dataset:

ID	0
date	0
Item Id	0
Item Name	344
ad_spend	1451
anarix_id	0
unit_price	0

dtype: int64







In [20]:

```
# Time series plot of units sold over time for a few items
sample_items = train['Item Id'].unique()[:5]
plt.figure(figsize=(14, 8))
for item in sample_items:
    item_data = train[train['Item Id'] == item]
    plt.plot(item_data['date'], item_data['units'], label=f'Item Id {item}')

plt.title('Units Sold Over Time for Sample Items')
plt.xlabel('Date')
plt.ylabel('Units Sold')
plt.legend()
plt.show()

# Average units sold by month
train['month'] = train['date'].dt.month
monthly_units = train.groupby('month')['units'].mean()
plt.figure(figsize=(10, 6))
monthly_units.plot(kind='bar')
plt.title('Average Units Sold by Month')
plt.xlabel('Month')
plt.ylabel('Average Units Sold')
plt.show()

# Ad spend vs units sold
plt.figure(figsize=(10, 6))
sns.scatterplot(data=train, x='ad_spend', y='units', alpha=0.5)
plt.title('Ad Spend vs Units Sold')
plt.xlabel('Ad Spend')
plt.ylabel('Units Sold')
plt.show()
```

```

# Correlation matrix
corr_matrix = train[['units', 'ad_spend', 'orderedrevenueamount', 'unit_price']].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Matrix')
plt.show()

# Top 10 items with highest average units sold
top_items = train.groupby('Item Id')['units'].mean().sort_values(ascending=False).head(10)
plt.figure(figsize=(12, 6))
top_items.plot(kind='bar')
plt.title('Top 10 Items with Highest Average Units Sold')
plt.xlabel('Item Id')
plt.ylabel('Average Units Sold')
plt.show()

```

```

In [15]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error
from statsmodels.tsa.arima.model import ARIMA

# Load the datasets
try:
    train = pd.read_csv('C:\\Users\\suhai\\Downloads\\train.csv\\train.csv')
    test = pd.read_csv('C:\\Users\\suhai\\Downloads\\test.csv')
    submission = pd.read_csv('C:\\Users\\suhai\\Downloads\\sample_submission.csv')
    print("Datasets loaded successfully.")
except Exception as e:
    print(f"Error loading datasets: {e}")

# Convert date columns to datetime
try:
    train['date'] = pd.to_datetime(train['date'])
    test['date'] = pd.to_datetime(test['date'])
    print("Date columns converted to datetime successfully.")
except KeyError as e:
    print(f"Error converting date columns: {e}")

```

```

# Feature Engineering
def create_features(df):
    df['month'] = df['date'].dt.month
    df['day'] = df['date'].dt.day
    df['weekday'] = df['date'].dt.weekday
    return df

try:
    train = create_features(train)
    test = create_features(test)
    print("Feature engineering completed successfully.")
except Exception as e:
    print(f"Error during feature engineering: {e}")

# Ensure data is sorted by date for each Item Id
try:
    train = train.sort_values(by=['Item Id', 'date'])
    print("Data sorted by 'Item Id' and 'date' successfully.")
except KeyError as e:
    print(f"Error sorting data: {e}")

# Model Training for each Item Id
results = []
for item in train['Item Id'].unique():
    item_data = train[train['Item Id'] == item].set_index('date')

```

```

# Ensure there are enough data points
if len(item_data) < 10:
    print(f"Skipping Item Id {item} due to insufficient data points")
    continue

try:
    # Define the model
    model = ARIMA(item_data['units'], order=(5, 1, 0))

    # Fit the model
    model_fit = model.fit()

    # Forecast for the test period
    start = len(item_data)
    end = start + len(test[test['Item Id'] == item]) - 1
    forecast = model_fit.predict(start=start, end=end, typ='levels')

    # Save the results
    forecast = pd.DataFrame(forecast, columns=['units'])
    forecast['date'] = test[test['Item Id'] == item]['date'].values
    forecast['Item Id'] = item
    results.append(forecast)
except Exception as e:
    print(f"Error fitting ARIMA model for Item Id {item}: {e}")

```

```

# Concatenate all results
try:
    results = pd.concat(results, ignore_index=True)
    print("Results concatenated successfully.")
except Exception as e:
    print(f"Error concatenating results: {e}")

# Debugging: Inspect the 'results' DataFrame
print("Results DataFrame columns:", results.columns)
print("Results DataFrame head:")
print(results.head())

# Prepare submission
# Ensure columns are correctly named before merging
if 'date' in results.columns and 'Item Id' in results.columns:
    try:
        submission = submission.merge(results, on=['date', 'Item Id'], how='left')
        submission = submission[['date', 'Item Id', 'units']]
        submission.to_csv('submission.csv', index=False)
        print("Forecasting and submission file creation complete.")
    except KeyError as e:
        print(f"Error preparing submission: {e}")
else:
    print("Error: 'date' or 'Item Id' column not found in results DataFrame.")

```

Output:

```
Skipping Item Id nan due to insufficient data points
Results concatenated successfully.
Results DataFrame columns: Index(['units', 'date', 'Item Id'], dtype='object')
Results DataFrame head:
   units  date      Item Id
0   NaN 2024-07-01  B09KDLQ2GW
1   NaN 2024-07-02  B09KDLQ2GW
2   NaN 2024-07-03  B09KDLQ2GW
3   NaN 2024-07-04  B09KDLQ2GW
4   NaN 2024-07-05  B09KDLQ2GW
```

NAP QUEENS ASSIGNMENT

Process Steps:

1. Import Libraries:

- Import pandas, numpy, matplotlib, statsmodels, and sklearn.

2. Load Datasets:

- Load the train, test, and submission datasets from the specified file paths.
- Handle any errors that occur during loading.

3. Convert Date Columns:

- Convert the date columns in the train and test datasets to datetime format.

4. Feature Engineering:

- Add new columns to both datasets: month, day, and weekday based on the date column.

5. Sort Data:

- Sort the training data by Item Id and date to ensure proper time series analysis.

6. Train ARIMA Model:

- For each unique Item Id in the training data:
 - Filter the data for that Item Id and set the date as the index.
 - Check if there are sufficient data points (at least 10).
 - Define and fit an ARIMA model on the units column.
 - Forecast sales for the test period based on the fitted model.
 - Collect the forecast results, including the date and Item Id.

7. Concatenate Results:

- Combine all individual forecasts into a single dataframe.

8. Prepare Submission:

- Merge the forecast results with the submission template.
- Ensure the final submission file contains date, Item Id, and units.
- Save the submission file as a CSV.

Insights and Results:

- The ARIMA model effectively captures the time series patterns of sales data for each item, generating forecasts based on historical trends.
- The feature engineering, including extraction of month, day, and weekday, provides additional context that can improve the model's accuracy.
- Sorting the data ensures that the model processes it chronologically, which is crucial for time series forecasting.
- Error handling during model fitting and data merging helps in identifying and resolving issues, ensuring reliable and accurate results.
- The final submission file integrates the forecasted units with the provided template, adhering to the competition's requirements and facilitating evaluation based on Mean Squared Error (MSE).

The process of forecasting sales using the ARIMA model involves detailed data preparation and feature engineering to handle temporal and categorical aspects of the data. By sorting the training data and applying ARIMA modeling to each item individually, the approach ensures that predictions are tailored to the sales patterns of each product. Despite facing challenges such as errors in merging data and issues with model fitting, the revised code successfully produces a forecast for each item and integrates these predictions with the submission format. The debugging steps and error handling improve the reliability of the process, allowing for accurate forecasting and generation of the submission file, which is critical for evaluating model performance. The final submission file is expected to reflect the forecasted sales, adhering to the competition's requirements.

